

Electric Vehicle Price Prediction in the Vietnamese Market

A Web-Scraped Dataset with LLM-Assisted Feature Extraction

Anh D. D. Hoang Khang H. Tran Anh N. Vu Nam N. Trinh

Information Technology Specialization
FPT University, Hanoi, Vietnam

AIL301m — March 2026



Outline

- 1 Introduction
- 2 Data Collection
- 3 Data Quality Analysis
- 4 Feature Engineering
- 5 Experimental Setup
- 6 Results
- 7 Discussion & Conclusion

- Vietnam's EV market is **growing rapidly**
 - VinFast mass-producing since 2022
 - Chinese imports (BYD, Wuling, Geely)
 - Government ICE phase-out by 2040
- **No structured dataset** exists for Vietnamese EV prices
- No local equivalent of Kelley Blue Book or AutoTrader
- Listings are **unstructured**, **Vietnamese-language**, scattered across multiple platforms

Gap

Existing car price prediction studies rely on clean, single-source, Western-market datasets

Our Goal

End-to-end pipeline: **scraping** → **LLM extraction** → **data quality** → **prediction**

Contributions

- ① **First Vietnamese EV dataset** — 4,166 records from 4 marketplaces with unified 20-column schema
- ② **Two-stage LLM extraction** — local Qwen 2.5 3B + cloud GPT-5-nano with auto-retry and Pydantic validation
- ③ **Data quality analysis** — identified and corrected 3 critical issues, reducing RMSE by **67%** (1.02B → 332M VND)
- ④ **Segment-specific evaluation** — practical accuracy ($\text{RMSE} < 100\text{M}$) for budget & mid-range segments (92% of test records)

Key Finding

Data quality remediation contributes **far more** to accuracy than model selection (67% vs 15% RMSE improvement).

Data Sources

Website	Method	Records
bonbanh.com	BeautifulSoup	1,493
chotot.com	Scrapy	1,666
xevinfastluot.com	Selenium	279
otodien.vn	BeautifulSoup	728
Total		4,166




Figure: Records per source

Unified schema: 20 columns covering ID, specs, condition, metadata, and price

Two-Stage Pipeline

① **Local — Qwen 2.5 3B** (via Ollama)

- Zero-shot prompt → JSON output
- Pydantic schema validation
- Auto-retry up to 3 attempts on failure

② **Cloud — GPT-5-nano** (verification)

- Batch-processes low-confidence records
- Corrects hallucinated trim names

Fill rates:

Brand: 99.3% Base model: 98.7%
Year: 98.3% Price: 100%

Failure mode:

Trim-level hallucination for non-VinFast brands

Three Critical Issues

1. chotot.com 10× Inflation

- 1,666 records (40%)
- chotot.com prices stored in 1,000 VND units
- Fix: divide by 10

2. BYD Pricing Error

- 90 records (2.3%)
- ~56% of total weighted error
- Fix: remove all BYD

3. Duplicate Leakage

- 887 duplicates (22%)
- Inflated R^2 : 0.787
- Honest R^2 : 0.725

Stage	Records	Best RMSE
Raw (all issues)	3,992	1.02B VND
+ All fixes applied	3,015	332M VND

↓ **67% RMSE reduction**

chotot.com Price Inflation




Figure: Systematic 10× price inflation from chotot.com vs other sources

Before vs After Data Quality Fixes




Figure: Model performance before and after corrections — “before” metrics were inflated by leakage

Feature Engineering Pipeline

41 features from 3,015 records:

Category	Count
Numeric (age, mileage, ...)	4
EV Specs (kWh, range, hp)	3
Target-encoded	2
One-hot (brand, model, ...)	32
Total	41

Key Decisions

- **EV specs recovery:** battery, range, power from otodien.vn lookup (29 models)
- **Hybrid encoding:** target-encode only high-cardinality features
- Target-encoding everything $\Rightarrow R^2$ drops to 0.087!

EV Specifications Recovery




Figure: EV specifications across models — captures fundamental pricing physics

Encoding Strategy & Feature Importance




Figure: Hybrid encoding




Figure: Top features by correlation & MI

Experimental Setup

Data Split:

- 80/20 train-test (2,412 / 603)
- Stratified on `base_model`
- Sample weights: inverse quintile frequency

Evaluation:

- RMSE, MAE, R^2 , Adj. R^2 , MAPE
- 95% CI via bootstrap (1,000 resamples)
- Segment-specific analysis

Four Models:

Model	Input
Ridge	Scaled + log target
SVR (RBF)	Scaled + log target
Random Forest	Raw features
XGBoost	Raw features

All tuned via **5-fold CV** with GridSearchCV

Overall Benchmark

Model	RMSE (M VND)	MAE (M VND)	R^2	Adj. R^2	MAPE (%)
Ridge	381	115	0.636	0.610	43.5
SVR	369	94	0.660	0.635	48.9
Random Forest	332	89	0.725	0.705	43.6
XGBoost	383	103	0.633	0.606	43.0

Table: Test set results ($n = 603$). RF 95% CI for RMSE: [139M, 496M] VND.

Random Forest wins across all metrics

Model Comparison




Figure: RMSE comparison




Figure: R^2 comparison

Segment-Specific Performance

Segment	RMSE	<i>n</i>
Budget (<500M)	100M	237
Mid (0.5–1.2B)	81M	318
Premium (1.2–5B)	797M	46
Luxury (>5B)	4,041M	2
Overall	332M	603

Table: RF by segment

Budget + Mid = **92% of test data** with practical RMSE






Figure: RMSE by segment (all models)

Random Forest: Predictions & Feature Importance



Data Quality vs Model Selection

Data Quality Fixes

RMSE: 1.02B → 332M VND

↓ **67% improvement**

- chotot.com price correction
- Deduplication (22% removed)
- BYD removal

Model Selection

Best vs Worst: 332M vs 383M

↓ **15% difference**

- RF vs XGBoost: 51M gap
- All models benefit equally from clean data

Takeaway

In applied ML with web-scraped data, **invest in data quality first** — it yields 4× more improvement than model experimentation.

Limitations:

- **Brand concentration** — VinFast = 93.6% of records
- **Small dataset** — 3,015 records (2,412 train)
- **No temporal modeling** — cross-sectional only
- **Luxury segment** — only 2 test records (n too small)

Future Work:

- ① Continuous scraping → time-series depreciation model
- ② Additional sources (dealer prices, registration data)
- ③ Ensemble stacking of all four models
- ④ Expand to hybrid & plug-in hybrid vehicles

Summary

- End-to-end pipeline: **scraping** → **LLM extraction** → **data quality** → **feature engineering** → **prediction**
- Data quality fixes reduced RMSE by **67%** (1.02B → 332M VND)
- Random Forest: $R^2 = 0.725$, RMSE = 332M VND
- **Practical accuracy** for budget + mid segments (92% of test records, RMSE $\leq$ 100M)

Data quality > Model selection

In real-world ML with scraped data, the dominant bottleneck is data quality, not model architecture.

Thank You

Questions?



FPT UNIVERSITY

Anh D.D. Hoang, Khang H. Tran, Anh N. Vu, Nam N. Trinh
FPT University, Hanoi — AIL301m, March 2026